


# Dynamic Programming

**Reinforcement Learning**  
October 13, 2022



# Learning Objectives

Understand the Bellman equation (better 😊 )

Distinguish between policy evaluation and control

Explain where dynamic programming can be used and where not

Explain how to compute value functions using policy iteration

Understand policy improvement

...

# Classical Dynamic Programming

Some basic ideas for the use of (classical) dynamic programming are:

- Divide and conquer
- Subdivide larger problems into smaller problems
- Solve smaller problems just once (and store solution)
- Make decisions in stages

## Example: 0-1 Knapsack Problem

Given a set of items with weights and values, pack a knapsack with given maximal weight to contain items of maximal total value

*# values and corresponding weights*

$v = [20, 5, 10, 40, 15, 25]$

$w = [1, 2, 3, 8, 7, 4]$

*# maximal weight*

$W = 10$



# Knapsack Problem

Solution:

- Recursively calculate a solution with/without the item
- Take the maximal value of both

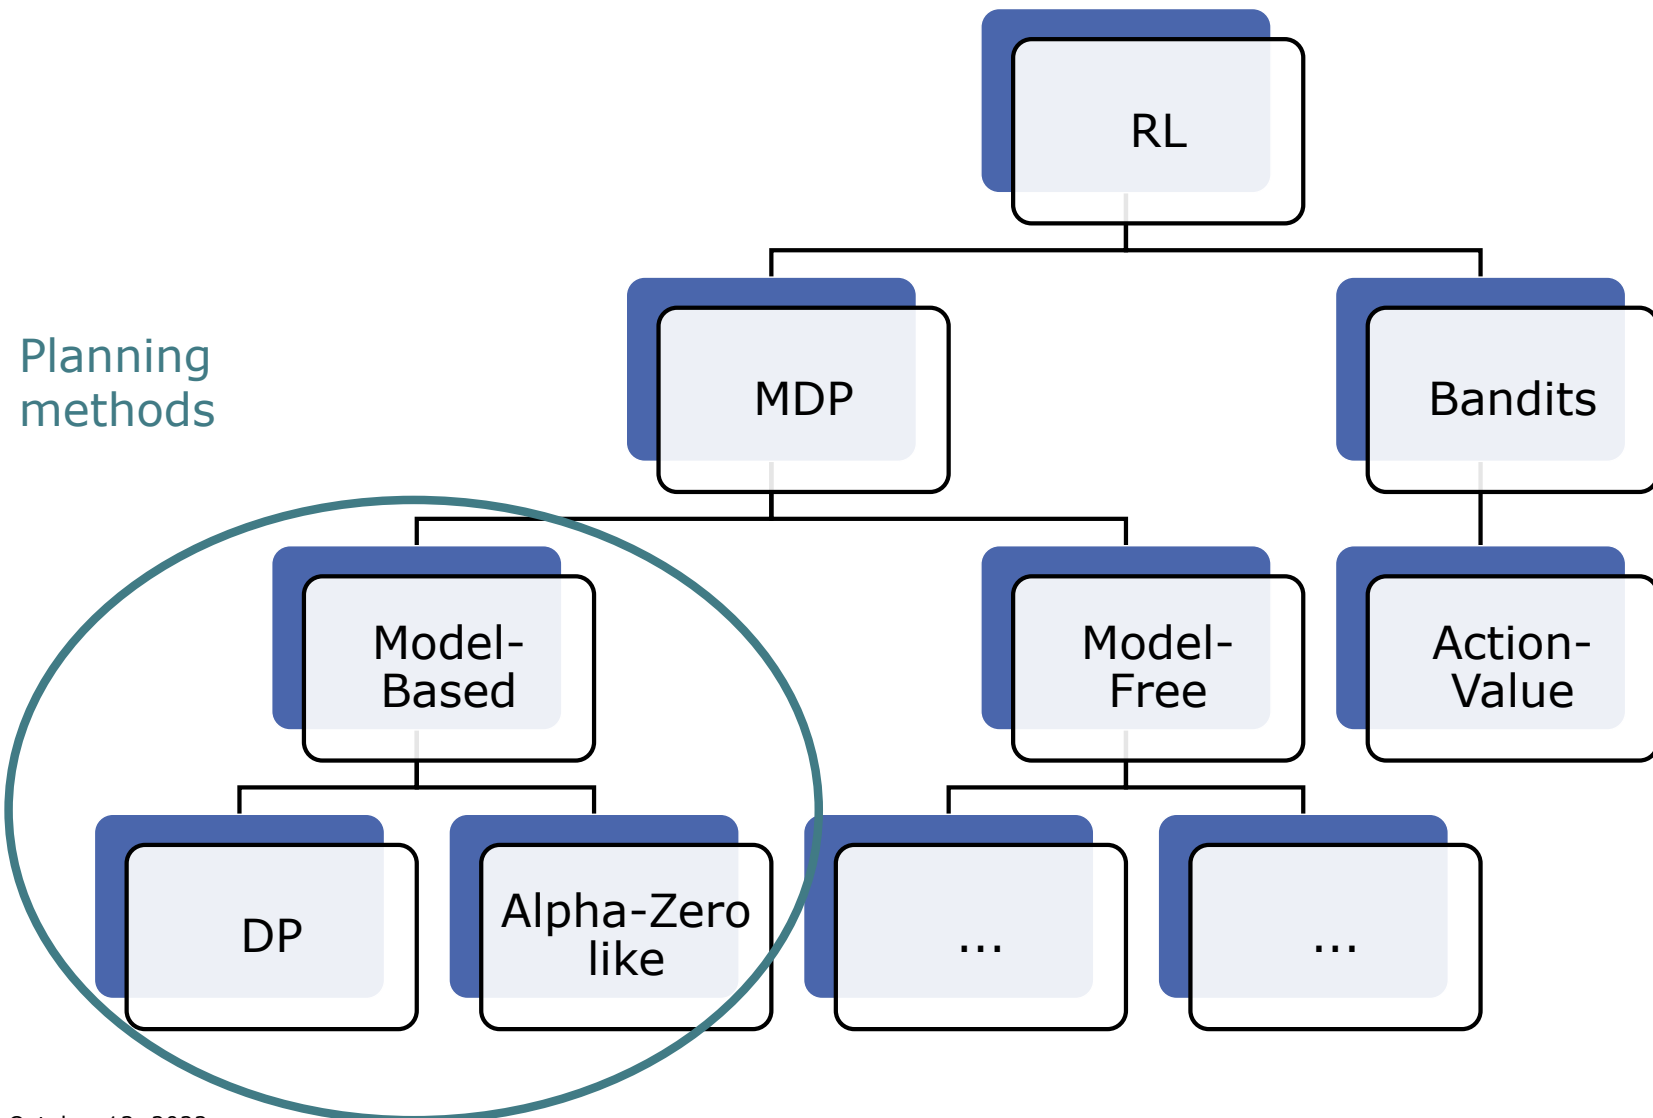
```
def knapsack(v, w, n, W):  
    """  
    v: values, w: weights, n: item to consider, W: weight left  
    """  
    if W < 0:  
        return -sys.maxsize  
    if n < 0 or W == 0:  
        return 0  
    include = v[n] + knapsack(v, w, n - 1, W - w[n])  
    exclude = knapsack(v, w, n - 1, W)  
    return max(include, exclude)
```

# Dynamic Programming in RL

Key ideas:

- Algorithms to compute optimal value functions and policies given a perfect **model** of the MDP
- Finite MDP environment
- Use value functions to organize the search for good policies
- An optimal policy can be obtained from an optimal value function
- DP is often limited in RL due to the need of a model and the large computational expense

# RL Methods



# Bellman Equation

$$v_{\pi}(s) = \sum_a \pi(a|s) \sum_{s',r} p(s',r | s,a) [r + \gamma v_{\pi}(s')], \text{ for all } s \in \mathcal{S}$$

- Calculates the value of a state by following policy  $\pi$
- Could be solved for  $v_{\pi}$  using a system of linear equations, but...
- ... iterative solutions are usually preferred (as they are computationally more efficient)



## Policy Evaluation (Prediction)

- Given a policy, what is its value function?
- Iterative computation using the Bellman equation
- Calculate a sequence of values that approach the correct solution

$$\begin{aligned} v_{k+1}(s) &\doteq \mathbb{E}_{\pi}[R_{t+1} + \gamma v_k(S_{t+1}) \mid S_t = s] \\ &= \sum_a \pi(a|s) \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_k(s')] \end{aligned}$$

- This is called *iterative policy evaluation*.

## Iterative Policy Evaluation, estimate $v_\pi$

Input: a policy  $\pi$

Initialize:

$V(s) \in \mathbb{R}$  arbitrarily, except  $V(\text{terminal}) = 0$

Loop:

$\Delta \leftarrow 0$

Loop for each  $s \in \mathcal{S}$ :

$v \leftarrow V(s)$


$V(s) \leftarrow \sum_a \pi(a|s) \sum_{s',r} p(s',r | s,a) [r + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until  $\Delta < \theta$

# Example for Policy Evaluation

- Small *Gridworld* Example
- Find terminal state (grey) from any position in minimal number of steps
- Start with random policy
- Initialize all values with 0 (any other value is possible too, except for terminal states which must be 0)



	1	2	3
4	5	6	7
8	9	10	11
12	13	14	

$R_t = -1$   
on all transitions

# Policy Evaluation

- The iterative policy evaluation calculates increasingly better estimates of the value function

- (Example from the book, values are given in only 1 digit precision)

$V_k$  for the  
Random Policy

$k = 0$

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

$k = 3$

0.0	-2.4	-2.9	-3.0
-2.4	-2.9	-3.0	-2.9
-2.9	-3.0	-2.9	-2.4
-3.0	-2.9	-2.4	0.0

$k = 1$

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0

$k = 10$

0.0	-6.1	-8.4	-9.0
-6.1	-7.7	-8.4	-8.4
-8.4	-8.4	-7.7	-6.1
-9.0	-8.4	-6.1	0.0

$k = 2$

0.0	-1.7	-2.0	-2.0
-1.7	-2.0	-2.0	-2.0
-2.0	-2.0	-2.0	-1.7
-2.0	-2.0	-1.7	0.0

$k = \infty$

0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0

# Policy Improvement

What happens if we change an action  $a$ , but follow  $\pi$  otherwise?

Recall that

$$q_{\pi}(s, a) = \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_{\pi}(s')]$$

If  $\pi$  and  $\pi'$  are deterministic policies with

$$q_{\pi}(s, \pi'(s)) \geq v_{\pi}(s), \quad \forall s \in \mathcal{S}$$

(i.e. we take the first action from  $\pi'$  then follow  $\pi$ ), then

$$v_{\pi'}(s) \geq v_{\pi}(s), \quad \forall s \in \mathcal{S}$$

# Policy Improvement

So, if we take a greedy action  $a$ , this defines a new policy

$$\begin{aligned}\pi'(s) &\doteq \operatorname{argmax}_a q_\pi(s, a) \\ &= \operatorname{argmax}_a \mathbb{E}[R_{t+1} + \gamma v_\pi(S_{t+1}) | S_t = s, A_t = a] \\ &= \operatorname{argmax}_a \sum_{s', r} p(s', r | s, a) [r + \gamma v_\pi(s')]\end{aligned}$$

which is as good as, or better than the old policy  $\pi$

# Policy Improvement

$V_k$  for the  
Random Policy

$k = 0$

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

$k = 1$

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0

$k = 2$

0.0	-1.7	-2.0	-2.0
-1.7	-2.0	-2.0	-2.0
-2.0	-2.0	-2.0	-1.7
-2.0	-2.0	-1.7	0.0

Greedy Policy  
w.r.t.  $V_k$

	↔	↔	↔
↔	↔	↔	↔
↔	↔	↔	↔
↔	↔	↔	

← random  
policy

	←	↔	↔
↑	↔	↔	↔
↔	↔	↔	↓
↔	↔	→	

	←	←	↔
↑	↖	↔	↓
↑	↔	↗	↓
↔	→	→	

# Policy Improvement

$k = 3$




0.0	-2.4	-2.9	-3.0
-2.4	-2.9	-3.0	-2.9
-2.9	-3.0	-2.9	-2.4
-3.0	-2.9	-2.4	0.0

$k = 10$

0.0	-6.1	-8.4	-9.0
-6.1	-7.7	-8.4	-8.4
-8.4	-8.4	-7.7	-6.1
-9.0	-8.4	-6.1	0.0

$k = \infty$

0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0




optimal  
policy



# Policy Iteration

Once a policy  $\pi$  has been improved using  $v_\pi$  to yield a better policy  $\pi'$ , we can compute  $v_{\pi'}$  and improve it again:

$$\pi_0 \xrightarrow{\text{E}} v_{\pi_0} \xrightarrow{\text{I}} \pi_1 \xrightarrow{\text{E}} v_{\pi_1} \xrightarrow{\text{I}} \pi_2 \xrightarrow{\text{E}} \dots \xrightarrow{\text{I}} \pi_* \xrightarrow{\text{E}} v_*$$



# Policy Iteration

Policy Iteration, estimate  $\pi \approx \pi_*$

Initialize:

$V(s) \in \mathbb{R}$  and  $\pi(s)$  arbitrarily, except  $V(\text{terminal}) = 0$

Loop:

Policy Evaluation:

estimate  $V \approx v_\pi$  (see previous algorithm)

Policy Improvement:

$\text{policy-stable} \leftarrow \text{true}$

For each  $s \in \mathcal{S}$ :

$\text{old-action} \leftarrow \pi(s)$

$\pi(s) \leftarrow \operatorname{argmax}_a \sum_{s'} p(s', r | s, a) [r + \gamma V(s')]$

if  $\text{old-action} \neq \pi(s)$ , then  $\text{policy-stable} \leftarrow \text{false}$

until  $\text{policy-stable}$

return  $V \approx v_*, \pi \approx \pi_*$

# Value Iteration

- Does the policy evaluation need to converge? (see gridworld example) ?
- If we stop policy evaluation after one sweep, we obtain an algorithm called *value iteration*
- The update for this can be written as

$$v_{k+1}(s) = \max_a \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_k(s')], \quad \forall s \in \mathcal{S}$$

# Value Iteration

Value Iteration, estimate  $\pi \approx \pi_*$

Initialize:

$V(s) \in \mathbb{R}$  arbitrarily, except  $V(\text{terminal}) = 0$

Loop:

$\Delta \leftarrow 0$

Loop for each  $s \in \mathcal{S}$ :

$v \leftarrow V(s)$

$V(s) \leftarrow \max_a \sum_{s',r} p(s',r \mid s,a) [r + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until  $\Delta < \theta$

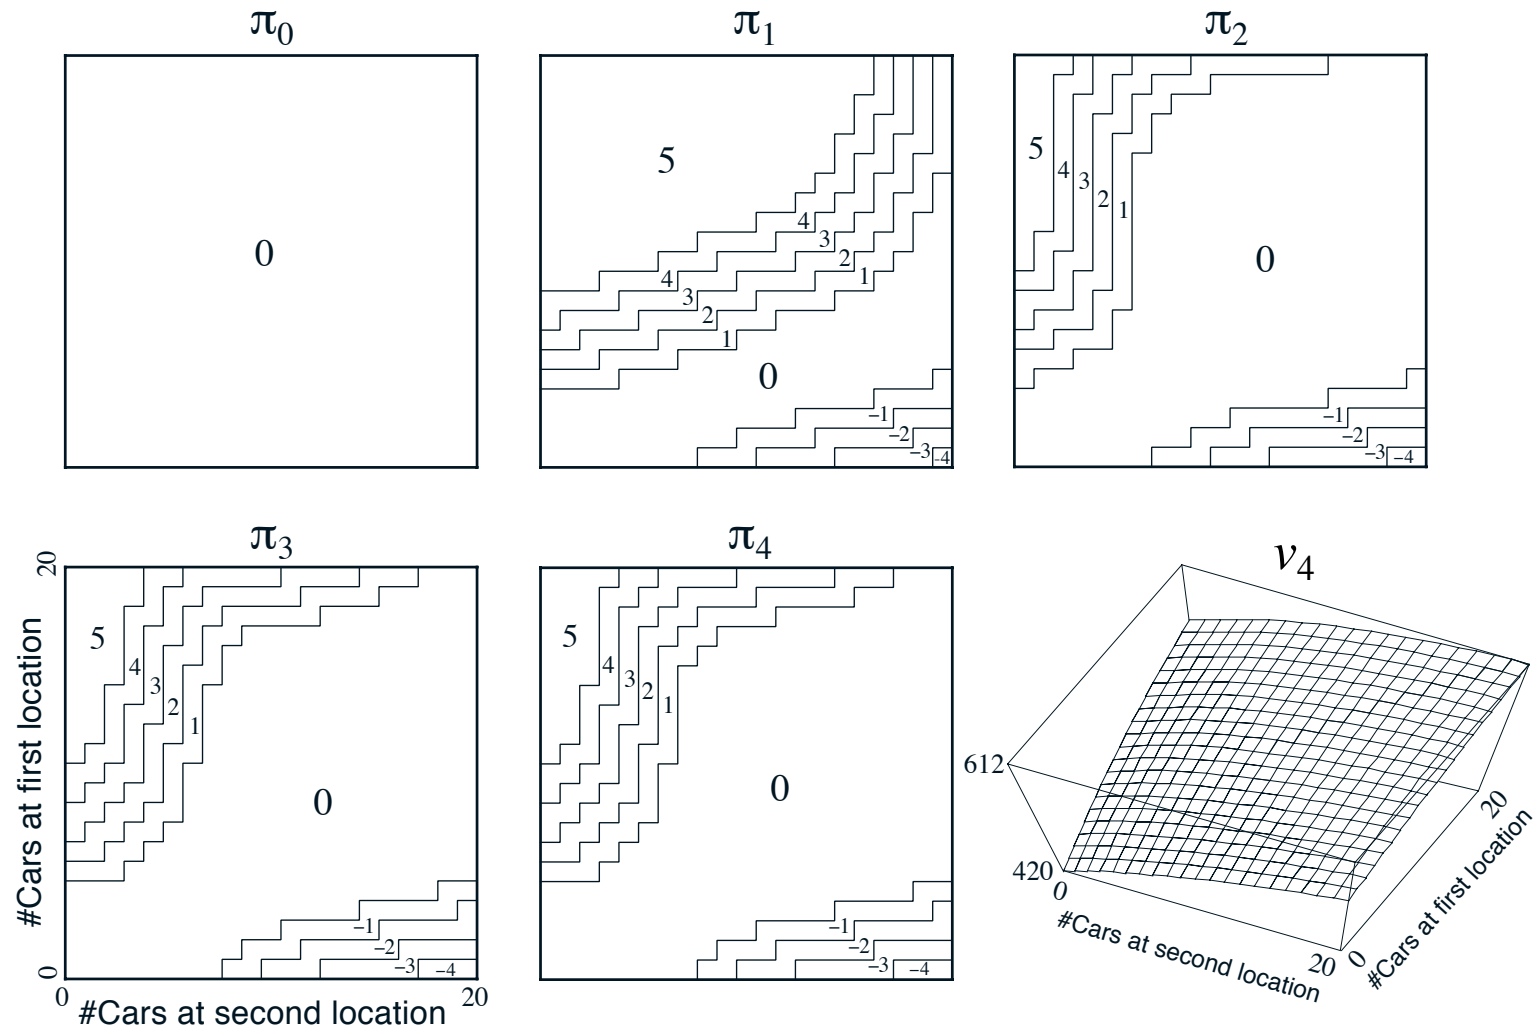
Output deterministic policy, such that

$\pi(s) = \operatorname{argmax}_a \sum_{s',r} p(s',r \mid s,a) [r + \gamma V(s')]$

## Example: Jack's Car Rental

- 2 car rental locations with max. 20 cars each
- Cars do not have to be returned to the same place
- If cars are available, they can be rented out for \$10, if a customer arrives
- Jack can move up to 5 cars during the night at a cost of \$2 (this is the action)
- (Poisson model of how the cars are rented and returned, not equal in both locations)
- Policy: Move  $n$  cars from first to second location
- What is the optimal policy? I.e., how many cars should be moved in each different state?

# Example: Jack's Car Rental



# Asynchronous vs. Synchronous Programming

## Synchronous DP:

- Update value function in sweeps
- All new values are calculated from the old values
- (Generally requires 2 arrays: one for the old values and one for the new ones)

## Asynchronous DP:

- Update value function in any order
- Update values in place
- All new values are calculated from the current values

# Generalized Policy Iteration

Generalized Policy Iteration (GPI):

- Interaction between *policy-evaluation* and *policy-improvement*
- (*independent of the granularity of the processes, i.e., if they are run just for one step or until convergence*)
- Many reinforcement learning methods can be described as GPI

